

In [1]:

```
#do not change
import sys
import os
from time import time
%matplotlib inline
from urllib.request import urlopen
import matplotlib.pyplot as plt
from sklearn.decomposition import PCA
import numpy as np
#from scipy.misc import imsave
import math
from tqdm import tqdm

def load_mnist():
    images_url = 'https://github.com/guptashvm/Data/blob/master/data/train-images-idx3-ubyte?raw=true'
    with urlopen(images_url) as urlopened:
        fd = urlopened.read()
        loaded = np.frombuffer(fd,dtype=np.uint8)
        trX = loaded[16:].reshape((60000,28*28)).astype(float)

    labels_url = 'https://github.com/guptashvm/Data/blob/master/data/train-labels-idx1-ubyte?raw=true'
    with urlopen(labels_url) as urlopened:
        fd = urlopened.read()
        loaded = np.frombuffer(fd,dtype=np.uint8)
        trY = loaded[8:].reshape((60000))

    trY = np.asarray(trY)

    X = trX / 255.
    y = trY

    subset = [i for i, t in enumerate(y) if t in [1, 0, 2, 3]]
    X, y = X.astype('float32')[subset], y[subset]
    return X[:1000], y[:1000]
```


Run the following code to load the data we need. Here, **X** is the array of images, **y** is the array of labels, and **X2d** is the array of projections of the images into two-dimensional space using PCA. The two-dimensional scatterplot of the top two principal components of the images is displayed, where the color of each point represents its label.

In [2]:

```
X, y = load_mnist()
pca = PCA(n_components=2)
pca.fit(X)
X2d = X.dot(pca.components_.T)

def plot_with_colors(Xs, ys):
    for i, _ in enumerate(ys):
        if ys[i] == 0:
            plt.plot([Xs[i, 0]], [Xs[i, 1]], 'r.')
        elif ys[i] == 1:
            plt.plot([Xs[i, 0]], [Xs[i, 1]], 'b.')
        elif ys[i] == 2:
            plt.plot([Xs[i, 0]], [Xs[i, 1]], 'g.')
        elif ys[i] == 3:
            plt.plot([Xs[i, 0]], [Xs[i, 1]], 'y.')
    plt.show()

plot_with_colors(X2d, y)
```



(a) Implement the standard k-means algorithm. Please complete the function **kmeans** defined below. You are NOT allowed to use any existing code of **kmeans** for this problem.

In [3]:

```
def kmeans(X, k = 4, max_iter = 500, random_state=0):
    """
    Inputs:a
        X: input data matrix, numpy array with shape (n * d), n: number of data points, d: feature dimension
        k: number of clusters
        max_iters: maximum iterations
    Output:
        clustering label for each data point
    """
    assert len(X) > k, 'illegal inputs'
    np.random.seed(random_state)

    # randomly select k data points as centers
    idx = np.random.choice(len(X), k, replace=False)
    centers = X[idx]

    # please complete the following code:

    from scipy.spatial import distance
    for i in range(max_iter):
        H = distance.cdist(X, centers, 'euclidean')
        labels = np.argmin(H, axis=1)
        centers = np.array([np.mean(X[labels==i], axis=0) for i in range(k)])

    return labels
```

(b) Run your **kmeans** function on the dataset (of the top two PCA components given by array X_{2d}). Set the number of clusters to 4. Visualize the result by coloring the 2D points in (a) according to their **clustering labels** returned by your **kmeans** algorithm. Because **kmeans** is sensitive to initialization, repeat your **kmeans** at least 5 times with different random initializations and show the plot of each initialization.

To quantitatively evaluate the clustering performance, we evaluate the accuracy , which can be written as follows, $\text{accuracy} = \max_{\mathcal{M}} \frac{\sum_{i=1}^n \mathbb{I}(y_i = \mathcal{M}(z_i))}{n}$, $n = 1000$, where y_i is the ground-truth label, z_i is the cluster assignment produced by the algorithm, and \mathcal{M} ranges over all possible one-to-one mapping between clusters and labels and $\mathbb{I}(x)$ is a indicator function ($\mathbb{I}(x) = 1 \text{ if } x=1; 0 \text{ otherwise}$). Please use the **accuracy_score** function defined below to calculate the accuracy. Report the best clustering accuracy you get out of 10 random initializations.

In [4]:


```
# do not change
def accuracy_score(y_true, y_pred):
    """
    Calculate clustering accuracy.
    y_true: true labels, numpy.array with shape `(n_samples,)`
    y_pred: predicted labels, numpy.array with shape `(n_samples,)`
    # Return
    accuracy, in [0,1]
    """
    y_true = np.squeeze(y_true)
    y_pred = np.squeeze(y_pred)
    assert y_true.shape == y_pred.shape, 'illegal inputs'

    y_true = y_true.astype(np.int64)
    assert y_pred.size == y_true.size
    D = max(y_pred.max(), y_true.max()) + 1
    w = np.zeros((D, D), dtype=np.int64)
    for i in range(y_pred.size):
        w[y_pred[i], y_true[i]] += 1
    from scipy.optimize import linear_sum_assignment
    row_ind, col_ind = linear_sum_assignment(w.max() - w)
    return sum([w[row_ind[i], col_ind[i]] for i in range(len(row_ind))]) * 1.0 /
y_pred.size
```


In [5]:

```
random_seeds = [1, 19, 1777, 51, 29, 314, 272, 2023, 13, 17, 1993]
for rs in random_seeds:
    labels = kmeans(X2d, random_state=rs, max_iter=10) # set max_iter=10 in order to see differences with
                                                       # different random initialization
    print(round(accuracy_score(labels, y)*100, 4))
    plot_with_colors(X2d, labels)
```


74.4




74.3




75.7



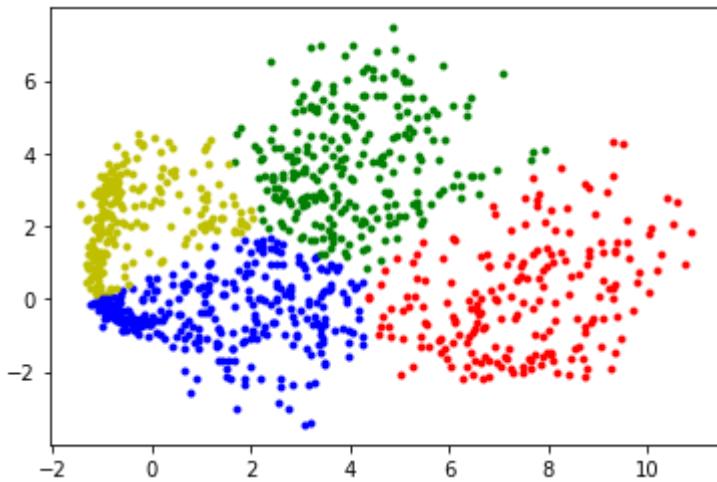
70.3




74.4




65.8




74.3




74.3




74.3



74.4



73.8



We have been testing k -means on the top two principal components for the purpose of visualization. Please run k -means on the (784 dimensional) original image dataset (again using 4 clusters). Try at least 10 different random initializations and report the best accuracy as above.

In [10]:

```
best_acc = -1
for rs in tqdm(random_seeds):
    labels = kmeans(X, random_state=rs)
    acc = accuracy_score(labels, y)
    if acc > best_acc:
        best_acc = acc
print(f"Best accuracy on the MNIST dataset using the k-means algorithm is {best_acc*100:.2f}")
```

100% |██████████| 11/11 [00:29<00:00, 2.70s/it]

Best accuracy on the MNIST dataset using the k-means algorithm is 8
6.10